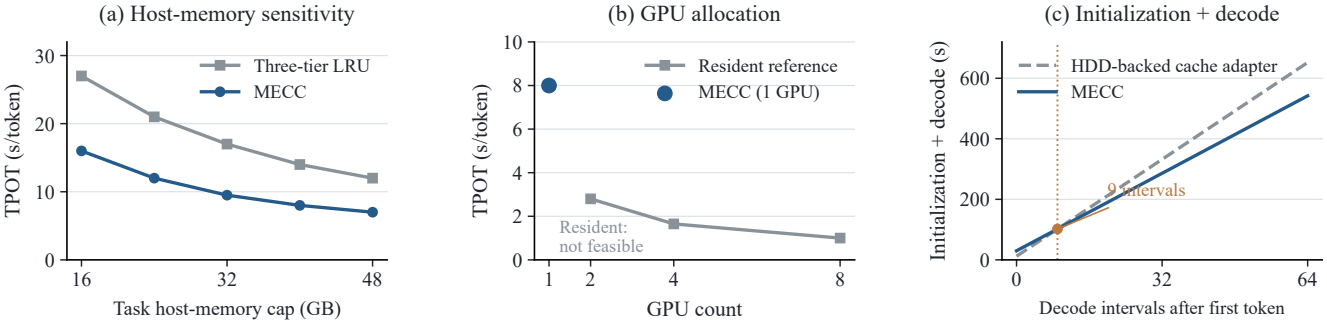
(a) Host-memory sensitivity
TPOT (s/token)
30
20
10
0
Three-tier LRU
MECC
16
32
48
Task host-memory cap (GB)
(b) GPU allocation
TPOT (s/token)
10
8
6
4
2
0
Resident reference
MECC (1 GPU)
Resident:
not feasible
1
2
4
8
GPU count
(c) Initialization + decode
Initialization + decode (s)
600
400
200
0
HDD-backed cache adapter
MECC
9 intervals
0
32
64
Decode intervals after first token