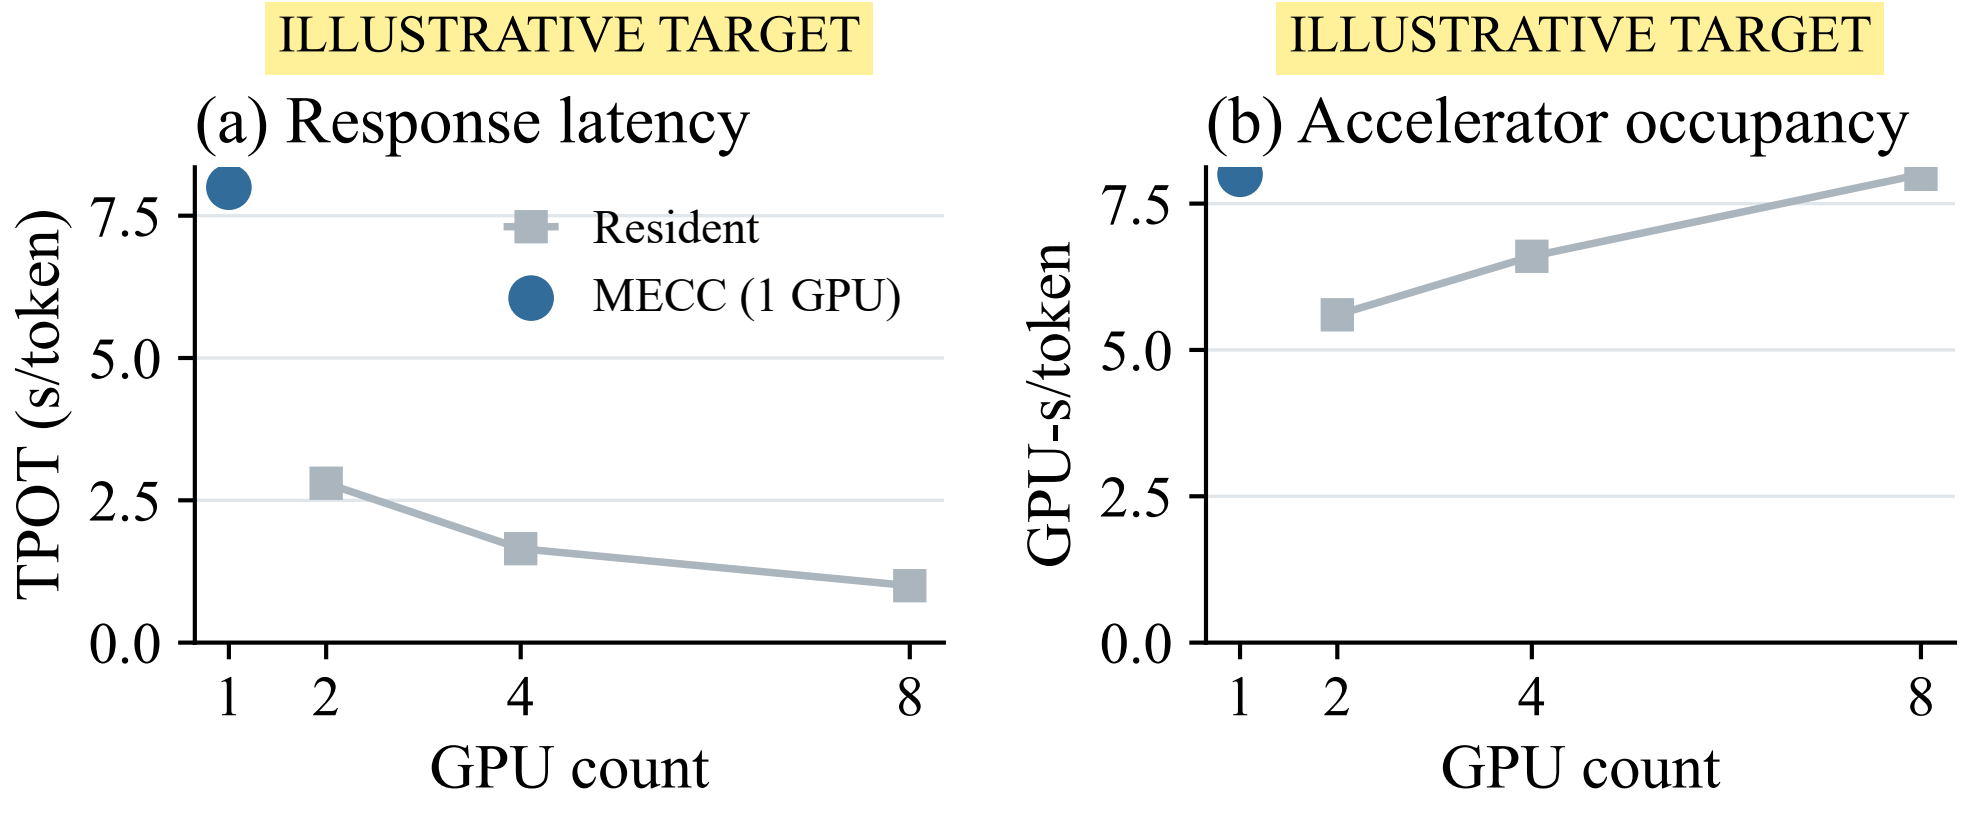
ILLUSTRATIVE TARGET
(a) Response latency
TPOT (s/token)
7.5
5.0
2.5
0.0
Resident
MECC (1 GPU)
1
2
4
8
GPU count
ILLUSTRATIVE TARGET
(b) Accelerator occupancy
GPU-s/token
7.5
5.0
2.5
0.0
1
2
4
8
GPU count